

به نام خدا

Explainable Security for Relational Databases

Sajad Asadi

Seyed Mohammad Mehdi Ahmadpanah

Computer Eng. And IT Department
Amirkabir University of Technology

May 29, 2016



دانشکده مهندسی کامپیوتر
و فناوری اطلاعات



دانشگاه صنعتی امیرکبیر
(پلی تکنیک تهران)

Reference

- ▶ Bender, Gabriel, Lucja Kot, and Johannes Gehrke. "Explainable security for relational databases." *Proceedings of the 2014 ACM SIGMOD international conference on Management of data*. ACM, 2014.

Outline

- ▶ Introduction
- ▶ Goal
- ▶ Database and Security Views
- ▶ Disclosure Control System
- ▶ Policy Function
- ▶ Disclosure Orders
- ▶ Disclosure Lattice
- ▶ Policy Formulas
- ▶ Explanations
- ▶ Disclosure Control In Practice

Goal

- ▶ A novel security model
- ▶ **Formal Explanation** for security decisions
- ▶ Query is accepted or denied? Why?

Database & Security Views

User			Friend	
UID	Name	Hobby	UID1	UID2
1	Babbage, Charles	math	1	3
2	Church, Alonzo	math	3	1
3	Lovelace, Ada	music	2	4
4	Turing, Alan	chess	4	2
5	von Neumann, John	history	4	5
			5	4

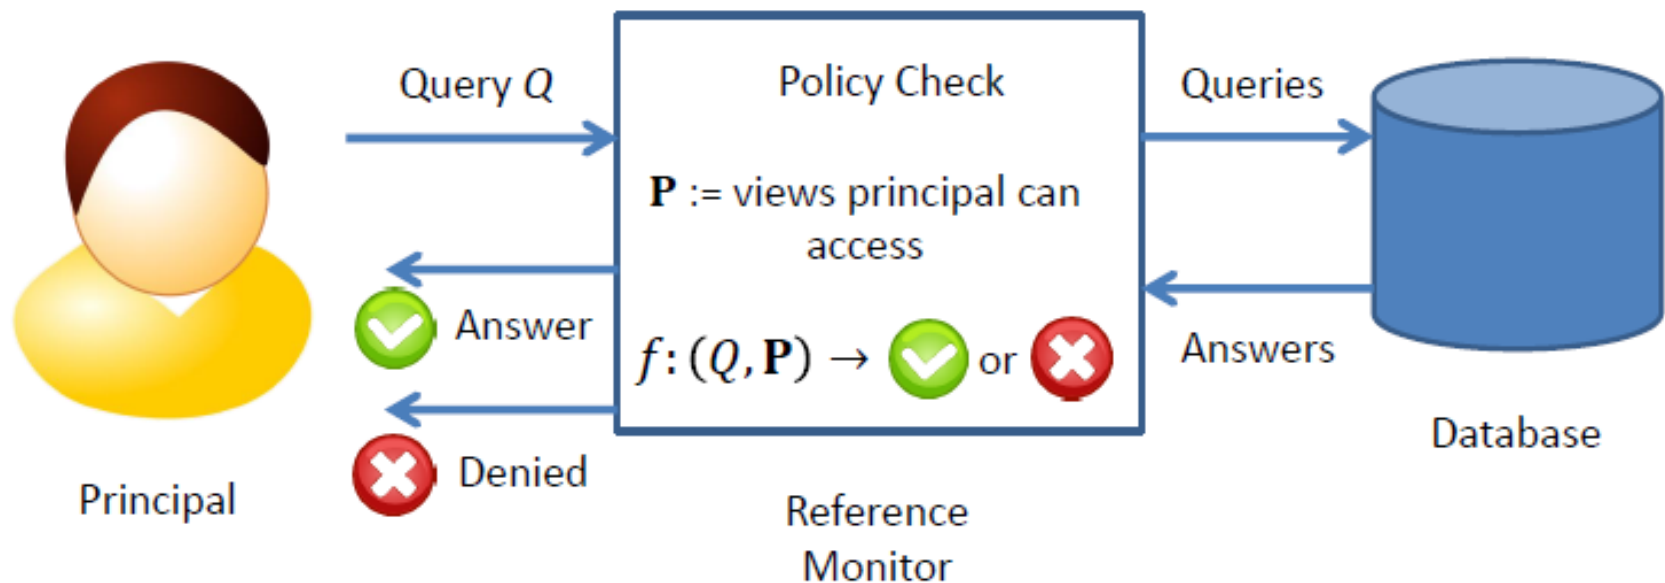
```
CREATE VIEW V1 AS  
(SELECT name, hobby FROM User WHERE uid = 1);
```

```
CREATE VIEW V2 AS  
(SELECT name FROM User WHERE uid = 1);
```

```
CREATE VIEW V3 AS  
(SELECT U.uid, U.name FROM User U, Friend F  
WHERE U.uid = F.uid2 AND F.uid1 = 1);
```

Written in Facebook Query Language

Disclosure Control System



Policy Function

- ▶ Variety of Policy Function:
 - Database-independent
 - Database-dependent
 - Not-data-derived
 - Data-derived

Example of not-data-derived

App has been granted to access V2 but not V1:

- ▶ `SELECT name FROM V1;`
- ▶ `SELECT name FROM V2;`

Two Different Semantics

In both relations and query answers:

- ▶ Set Semantics : No duplicates
- ▶ Bag Semantics : Can contain duplicate

Query Language:

Conjunctive Query under bag Semantics

- ▶ $H :- B$
- ▶ H (Head of query) is a relational atom.
- ▶ B (Body of query) is conjunction of relational atoms.
- ▶ Atom is constant or variable.

Example

```
CREATE VIEW V1 AS  
  (SELECT name, hobby FROM User WHERE uid = 1);
```

```
CREATE VIEW V2 AS  
  (SELECT name FROM User WHERE uid = 1);
```

```
CREATE VIEW V3 AS  
  (SELECT U.uid, U.name FROM User U, Friend F  
    WHERE U.uid = F.uid2 AND F.uid1 = 1);
```

$$V_1(n, h) :- U(1, n, h)$$
$$V_2(n) :- U(1, n, h)$$
$$V_3(u, n) :- U(u, n, h), F(1, u)$$

Disclosure Orders

- ▶ Some set of views (V') reveal more information about data set than others (V).

DEFINITION 1. A *disclosure order* is a binary relation on subsets of \mathcal{U} such that

- (i) If $V \subseteq V'$ then $V \preceq V'$.
- (ii) If $V \preceq V'$ and $V' \preceq V''$ then $V \preceq V''$.
- (iii) If $V \preceq V''$ and $V' \preceq V''$ then $V \cup V' \preceq V''$.

Example

- (i) **Set containment:** $V \preceq V'$ if and only if $V \subseteq V'$.
- (ii) **View determinacy:** $V \preceq V'$ if and only if the answers to the views in V' uniquely determine the answers to the views in V on every possible dataset.
- (iii) **View rewriting:** $V \preceq V'$ if and only if every view in V has a rewriting using V' .

Disclosure Orders to Disclosure Lattice Extension

We would like sets of views to form a lattice where the *greatest lower bound* (GLB) of \mathbf{V} and \mathbf{V}' (denoted $\mathbf{V} \sqcap \mathbf{V}'$) represents the information that is common to the two sets and the *least upper bound* (LUB) of \mathbf{V} and \mathbf{V}' (denoted $\mathbf{V} \sqcup \mathbf{V}'$) represents the combined information from the two sets.

Lattice Property

- ▶ Reflexivity
- ▶ Transitivity
- ▶ Not Antisymmetric!!!

$$(\Downarrow \mathbf{V}) = \{V \in \mathcal{U} : \{V\} \preceq \mathbf{V}\}$$

Theorem

THEOREM 1. Let \mathcal{U} be a set of views, and let \preceq be a disclosure order for \mathcal{U} . Define $\mathcal{I} = \{\Downarrow \mathbf{V} : \mathbf{V} \subseteq \mathcal{U}\}$. Then (\mathcal{I}, \preceq) is a lattice; details are as follows:

- LUB: $(\Downarrow \mathbf{V}) \sqcup (\Downarrow \mathbf{V}') = \Downarrow (\mathbf{V} \cup \mathbf{V}')$.
- GLB: $(\Downarrow \mathbf{V}) \sqcap (\Downarrow \mathbf{V}') = (\Downarrow \mathbf{V}) \cap (\Downarrow \mathbf{V}')$.
- Top element $\top = (\Downarrow \mathcal{U}) = \mathcal{U}$, bottom element $\perp = (\Downarrow \emptyset)$.

Policy Function

- ▶ Data-derived policy functions can be expressed in terms of disclosure lattices.

$$f(Q; \mathbf{P}) = 1 \text{ if and only if } \{Q\} \leq \mathbf{P}$$

Where \mathbf{P} is subset of security views.

$$\Downarrow \{Q\} \preceq \Downarrow \mathbf{P}$$

Policy Formulas

- ▶ Each view in V roughly corresponds to a single permission that a principal may or may not be granted.

$$\tau ::= 0 \mid 1 \mid V \mid (\tau \vee \tau) \mid (\tau \wedge \tau)$$

- ▶ $\mathbf{P} \vdash \varphi$ if and only if $f(Q;\mathbf{P}) = 1$

Means that φ is provable from \mathbf{P}

Examples

- ▶ $\varphi = V1$

- ▶ $\varphi = V1 \vee V2$

- ▶ $\varphi = V1 \wedge V3$

- ▶ $V2 \vdash V1 \vee V2$

What φ tells?

- ▶ φ tells us precisely which combinations of the security views contain enough information to uniquely determine the answer to Q .

Rules of Reference Monitor

$\mathbf{P} \vdash 1$ and $\mathbf{P} \not\vdash 0$

$\mathbf{P} \vdash V$ if and only if $V \in \mathbf{P}$

$\mathbf{P} \vdash (\varphi \vee \psi)$ if and only if $\mathbf{P} \vdash \varphi$ or $\mathbf{P} \vdash \psi$

$\mathbf{P} \vdash (\varphi \wedge \psi)$ if and only if $\mathbf{P} \vdash \varphi$ and $\mathbf{P} \vdash \psi$

Generating Data-Derived Formulas

```
1: procedure POLICY(Q, V)
2:    $\varphi \leftarrow 1$ 
3:   for  $Q \in \mathbf{Q}$  do
4:      $\psi \leftarrow 0$ 
5:     for  $V \in \mathbf{V}$  do
6:       if  $\{Q\} \preceq \{V\}$  then
7:          $\psi \leftarrow (\psi \vee V)$ 
8:       end if
9:     end for
10:     $\varphi \leftarrow (\varphi \wedge \psi)$ 
11:  end for
12:  return  $\varphi$ 
13: end procedure
```

Generating Data-Derived Formulas

$$\text{policy}(\mathbf{Q}) = \bigwedge_{Q \in \mathbf{Q}} \left(\bigvee_{V \in \mathbf{V}: \{Q\} \preceq \{V\}} V \right)$$

Example

For example, let \mathbf{Q} contain the queries

$$Q_{16}(1, n) :- U(1, n, h)$$

$$Q_{17}(h) :- U(u, n, h)$$

and let \mathbf{V} consist of the views from Figure 3. Then

$$\text{policy}(\{Q_{16}\}) = V_9 \vee V_{10} \vee V_{11}$$

$$\text{policy}(\{Q_{17}\}) = V_9 \vee V_{12}$$

and therefore

$$\text{policy}(\{Q_{16}, Q_{17}\}) = (V_9 \vee V_{10} \vee V_{11}) \wedge (V_9 \vee V_{12})$$

Explanations

- ▶ Why-so: 0 for every view in policy formula that principal has not been granted access to.
- ▶ Why-not: 1 for every view in policy formula that principal has been granted access to.

Example

$$\text{policy}(\{Q_{16}, Q_{17}\}) = (V_9 \vee V_{10} \vee V_{11}) \wedge (V_9 \vee V_{12})$$

Why-so:

$$(V_9 \vee V_{10} \vee 0) \wedge (V_9 \vee 0) = (V_9 \vee V_{10}) \wedge (V_9) = V_9$$

Why-not:

$$(V_9 \vee 1 \vee 1) \wedge (V_9 \vee V_{12}) = 1 \wedge (V_9 \vee V_{12}) = (V_9 \vee V_{12})$$

Disclosure Control In Practice

▶ Reflective Policies

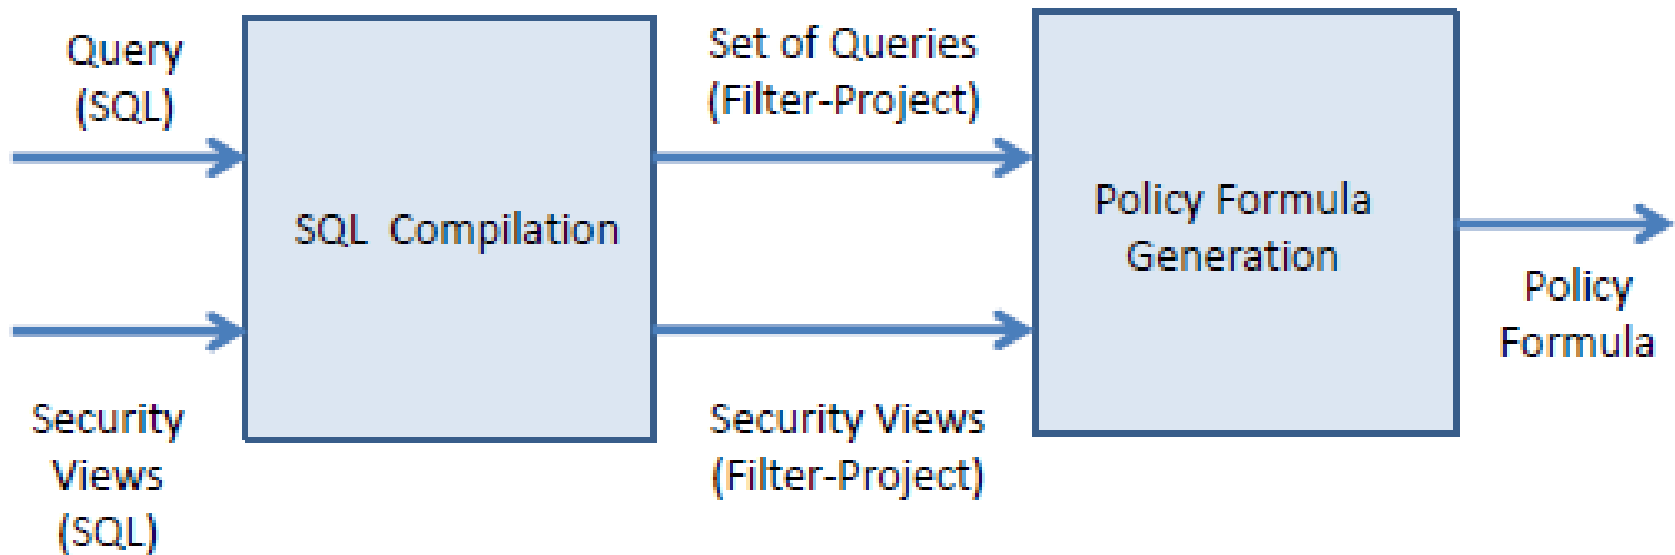
- policies governing tuples in one part of database can depend on a different part of the same database
- Frequently in Practice
 - e.g. Facebook's security policies

Running Example

```
SELECT U1.uid, U1.name FROM User U1 WHERE  
U1.uid IN  
(SELECT F1.uid2 FROM Friend F1 WHERE  
F1.uid1 = 1)
```

User			Friend	
UID	Name	Hobby	UID1	UID2
1	Babbage, Charles	math	1	3
2	Church, Alonzo	math	3	1
3	Lovelace, Ada	music	2	4
4	Turing, Alan	chess	4	2
5	von Neumann, John	history	4	5
			5	4

System Architecture



System Architecture (cont.)

- ▶ Each Security View: Filter-Project Query
 - SELECT-FROM-WHERE query with semi-joins but not inner or outer joins

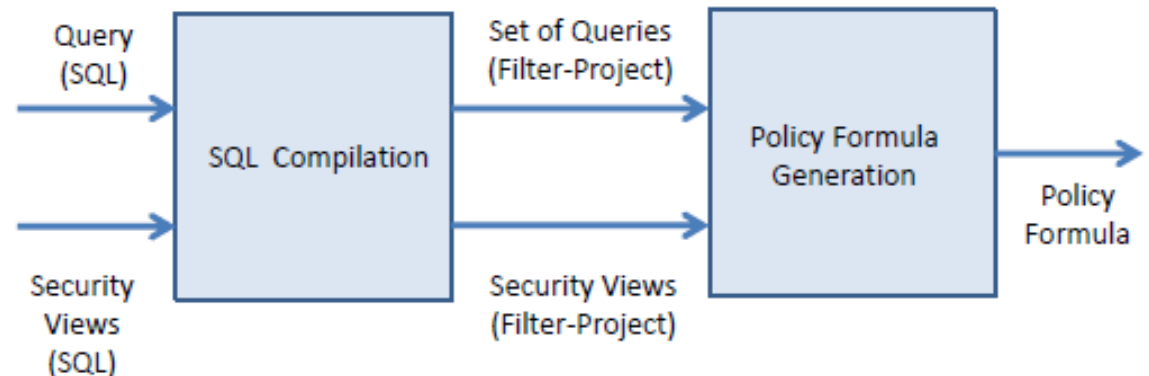
<i>Employee</i>		
Name	Empld	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Production

<i>Dept</i>	
DeptName	Manager
Sales	Bob
Sales	Thomas
Production	Katie
Production	Mark

<i>Employee ⋈ Dept</i>		
Name	Empld	DeptName
Sally	2241	Sales
Harriet	2202	Production

System Architecture (cont.)

- ▶ Compilation is conservative
 - Policy Formula → Upper bound on the information needed to answer a query
- ▶ Database-independent
 - Policy formula and Explanations



Filter–Project Queries

▶ Tables vs. Table Instances

- Query-independent vs. FROM clauses

```
SELECT U1.uid, U1.name FROM User U1, Friend F1  
WHERE F1.uid1 = 1 AND F1.uid2 = U1.uid
```

- ▶ Which *columns* and *rows* in each table instance can affect the query's output?

Filter–Project Queries (cont.)

$$H :- B \bowtie F_1 \wedge F_2 \wedge \dots$$

- ▶ A head atom H
- ▶ A main body atom B
- ▶ zero or more filter atoms F_1, F_2, \dots

Filter–Project Queries (cont.)

▶ Running Example:

SELECT U F1
 WHERE F1 $Q_{18}(u, n) :- U(u, n, h) \bowtie F(1, u)$

SELECT U1.uid, U1.name FROM User U1 WHERE U1.uid IN
 (SELECT uid2 FROM Friend WHERE uid1 = 1)

$Q_{19}(u) :- F(1, u) \bowtie U(u, n, h)$

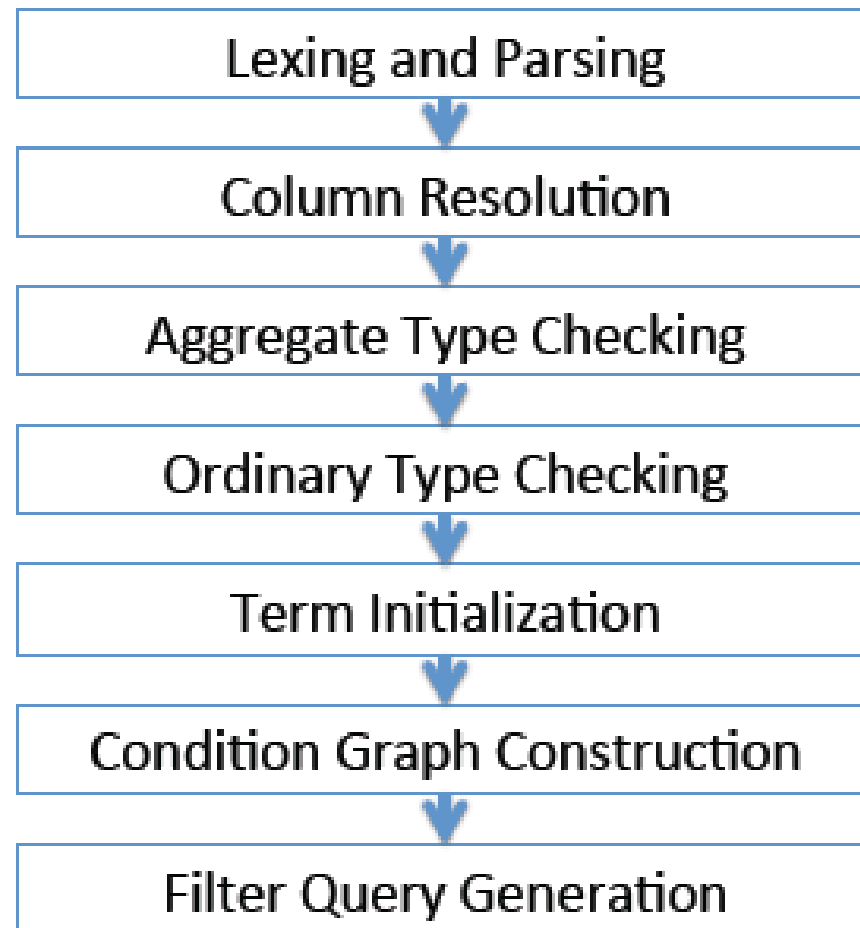
SELECT F1.uid2 FROM Friend F1 WHERE F1.uid1 = 1
 AND F1.uid2 IN (SELECT uid FROM User)

$Q_{19}(u) :- F(1, u) \bowtie U(u, n, n)$

Filter–Project Queries (cont.)

- ▶ Filter atoms in the query language correspond to semi-joins in standard SQL
- ▶ Order filter-project queries according to the amount of information they disclose about their *main body atoms*
 - RM will perform exactly one policy check for each table instance in the original SQL query
 - The universe of filter-project queries is *decomposable*

Query Compilation Pipeline



Preprocessing and Column Resolution

- ▶ Lexing and Parsing using JSqlParser
- ▶ Execution a series of passes on the *Abstract Syntax Tree* (ATS)
- ▶ Identify table instances
- ▶ Preform column resolution

Type Checking

- ▶ Aggressive type-checking on the AST

- *Aggregate type-checking*

SELECT SUM(SUM(U1.uid)) FROM User U1

- *Ordinary type-checking*

SELECT * FROM User U1 WHERE U1.uid IN
(SELECT F1.uid1, F1.uid2 FROM Friend F1)

Term Initialization

- ▶ Convert the AST for a SQL query Q into a collection of projection queries
 - Together reveal enough information to uniquely determine the answer to Q on any possible dataset
- 1. Generate one query for each table instance in Q
- 2. Determine the *existential* and *distinguished* variables in query

$$Q_{U1}(u, n) :- U(u, n, h)$$

$$Q_{F1}(u_1, u_2) :- F(u_1, u_2)$$


Generating Filter–Project Queries

3. Condition Graph


- Determine equality constraints and filter atoms to be added to the queries of the previous step
 - Reachability in the condition graph
- Similar to expression trees in the Relational Algebra

Generating Filter-Project Queries (cont.)


► Condition Graph




(a) Two-way join




(b) Left outer join



(c) Full outer join




(d) Three-way join



(e) Nested subquery

Example

```
SELECT U1.uid, U1.name FROM User U1 WHERE U1.uid IN  
(SELECT F1.uid2 FROM Friend F1 WHERE F1.uid1 = 1)
```



$$Q_{U1}(u, n) \quad :- \quad U(u, n, h) \bowtie F(u_1, u_2)$$
$$Q_{F1}(u_1, u_2) \quad :- \quad F(u_1, u_2) \bowtie U(u, n, h)$$

Generating Filter–Project Queries (cont.)

4. Find equality constraints to be added to the previous queries

- Finding all SELECCT nodes that are reachable from the condition graph node associated with the table instances
 - Then check WHERE and HAVING clauses

Example


$$\begin{aligned}Q_{U1}(u, n) &:- U(u, n, h) \bowtie F(u_1, u_2) \\Q_{F1}(u_1, u_2) &:- F(u_1, u_2) \bowtie U(u, n, h)\end{aligned}$$

Constraints: $(F1.uid1 = 1)$ and $(F1.uid2 = U1.uid)$

$$\begin{aligned}Q_{U1}(u, n) &:- U(u, n, h) \bowtie F(1, u) \\Q_{F1}(1, u_2) &:- F(1, u_2) \bowtie U(u_2, n, h)\end{aligned}$$

Experimental Evaluation

- Performance for ordinary queries
- Nearly linearly with the number of cores



Any Questions?



